

Foundations of Embedded Systems

Topic 04: Low-Level C Programming and Tools for Embedded Systems

(Video)

(~60 minutes)

Version 0.2020

Phillip Stanley-Marbell

Department of Engineering, University of Cambridge

<http://physcomp.eng.cam.ac.uk>

Intended Learning Outcomes for Today

2

36

By the end of this topic, you should be able to:

- ❶ Enumerate **steps in compiling a program** from C to machine instructions
- ❷ Enumerate and understand the **tools involved in the process**
- ❸ Enumerate the **main components of a processor (core) implementation**
- ❹ Understand and enumerate **how assembly code gets executed** in processor
- ❺ **Design** your own **linker script files** and use tools like **objdump** and **objcopy**

Next:

- ▶ We will look at **four different variants of a simple 4-line C program** to introduce the compilation and linking process

Anatomy of Four C Programs: Program 1

4

35

```
1 int
2 main(void) {
3     return 0;
4 }
```

```
sh-coff-gcc -E init.S > init.i; sh-coff-as init.i -o init.o
```

```
sh-coff-gcc -DM32 -Wall -gstabs3 -O0 -c simple.c
```

```
sh-coff-ld -Ttext 0x08004000 -TsuperH.ld -L../../../tools/tools-lib/superH -L
    ../../../sys/libOS/mOS -Map simple.map init.o simple.o -o simple -lc -lgcc -
    lm -lmOS-superH
```

```
sh-coff-objcopy -O srec simple simple.sr
```

More on **linker script files** (an *input*) later

More on **map files** (an *output*) later

Anatomy of Four C Programs: Disassembling Program 1

5

35

```
1 % sh-coff-objdump -d simple
2
3 simple:      file format coff-sh
4
5 Disassembly of section .text:
6
7 00000000008004000 <_text>:
8   8004000:    c9 00                and #0,r0
9   8004002:    40 0e                ldc r0,sr
10
11 ...      150 lines deleted      ...
12
13 80040ee:    00 09                nop
```

Disassembled binary is 160 lines

F(E)




Anatomy of Four C Programs: Program 2

6

35

```
1 int
2 noMain(void) {
3     return 0;
4 }
```



```
sh-coff-gcc -E init.S > init.i; sh-coff-as init.i -o init.o
sh-coff-gcc -DM32 -Wall -gstabs3 -O0 -c simple.c
sh-coff-ld -Ttext 0x08004000 -TsuperH.ld -L../../../tools/tools-lib/superH -L
    ../../../sys/libOS/mOS -Map simple.map init.o simple.o -o simple -lc -lgcc -
    lm -lmOS-superH
sh-coff-objcopy -O srec simple simple.sr
```

Exactly the same compilation flags and steps

F(E)

(10)

Anatomy of Four C Programs: Disassembling Program 2

7

35

```
1 % sh-coff-objdump -d simple
```

2

```
3 simple:      file format coff-sh
```

4

```
5 Disassembly of section .text:
```

6

```
7 00000000008004000 <_text>:
```

```
8 8004000:      c9 00          and #0,r0
```

```
9 8004002:      40 0e          ldc r0,sr
```

```
10 8004004:      df 02          mov.l    8004010 <stack_addr>,r15      ! 0xc001000
```

```
11 8004006:      d0 03          mov.l    8004014 <start_addr>,r0 ! 0x8004020 <_noMain>
```

```
12 8004008:      40 0b          jsr @r0
```

```
13 800400a:      00 09          nop
```

```
14 800400c:      e4 01          mov #1,r4
```

```
15 800400e:      c3 22          trapa   #34
```

What will be the difference in binaries between Program 1 and Program 2?


Disassembled binary is just 38 lines. Only difference is change of function name

Anatomy of Four C Programs: Program 3

8


35

```
1 #include <stdio.h>
2
3 int
4 main(void) {
5     return 0;
6 }
```



Compared to Program 1:

```
1 int
2 main(void) {
3     return 0;
4 }
```



What will be the difference in binaries between Program 1 and Program 3?



Absolutely no difference

Anatomy of Four C Programs: Program 4

9



35

```
1 #include <stdio.h>
2
3 int
4 main(void) {
5     printf("Hello, World\n");
6
7     return 0;
8 }
```



(Compared to Program 3:)

```
1 #include <stdio.h>
2
3 int
4 main(void) {
5     return 0;
6 }
```



Disassembled binary will now be 4116 lines (vs. 160 lines for Progs. 1 and 3)

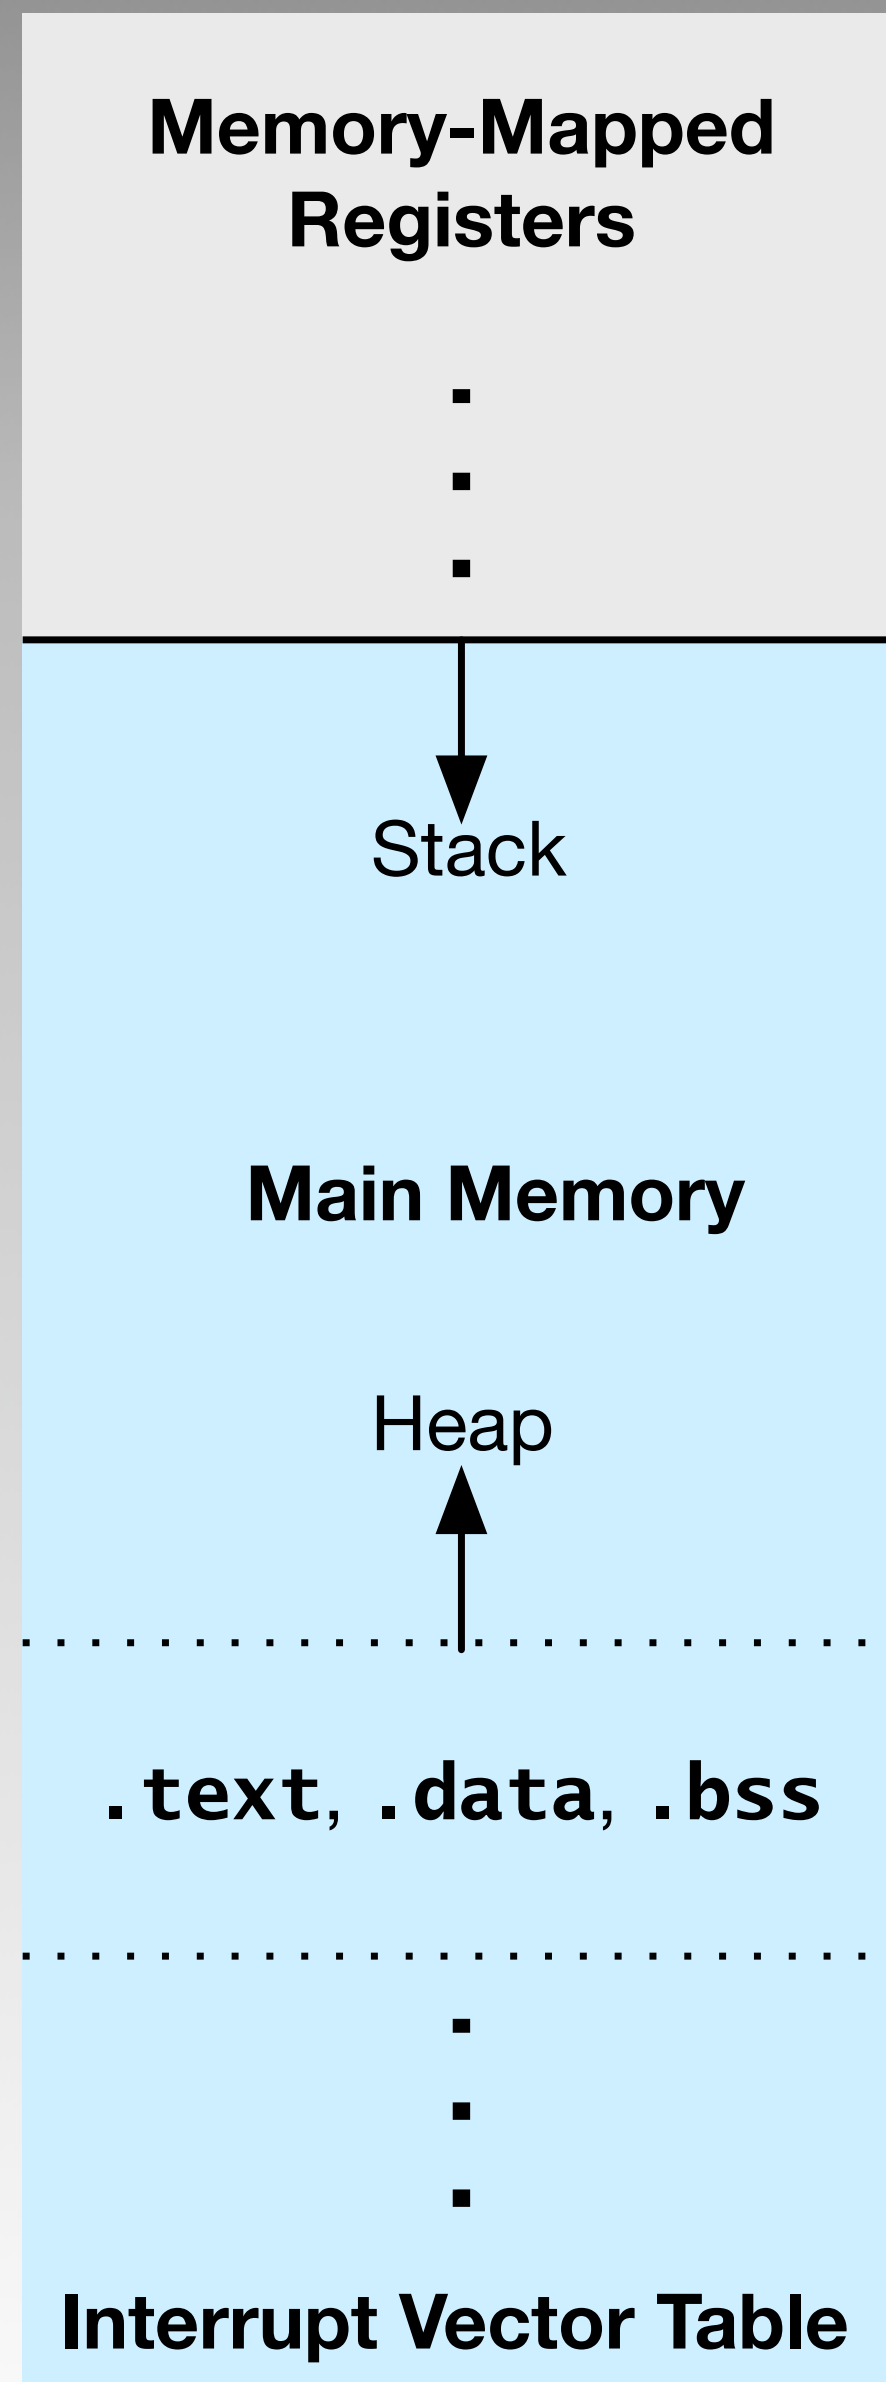


Memory Map and the Parts of a Binary: Terminology

10
35

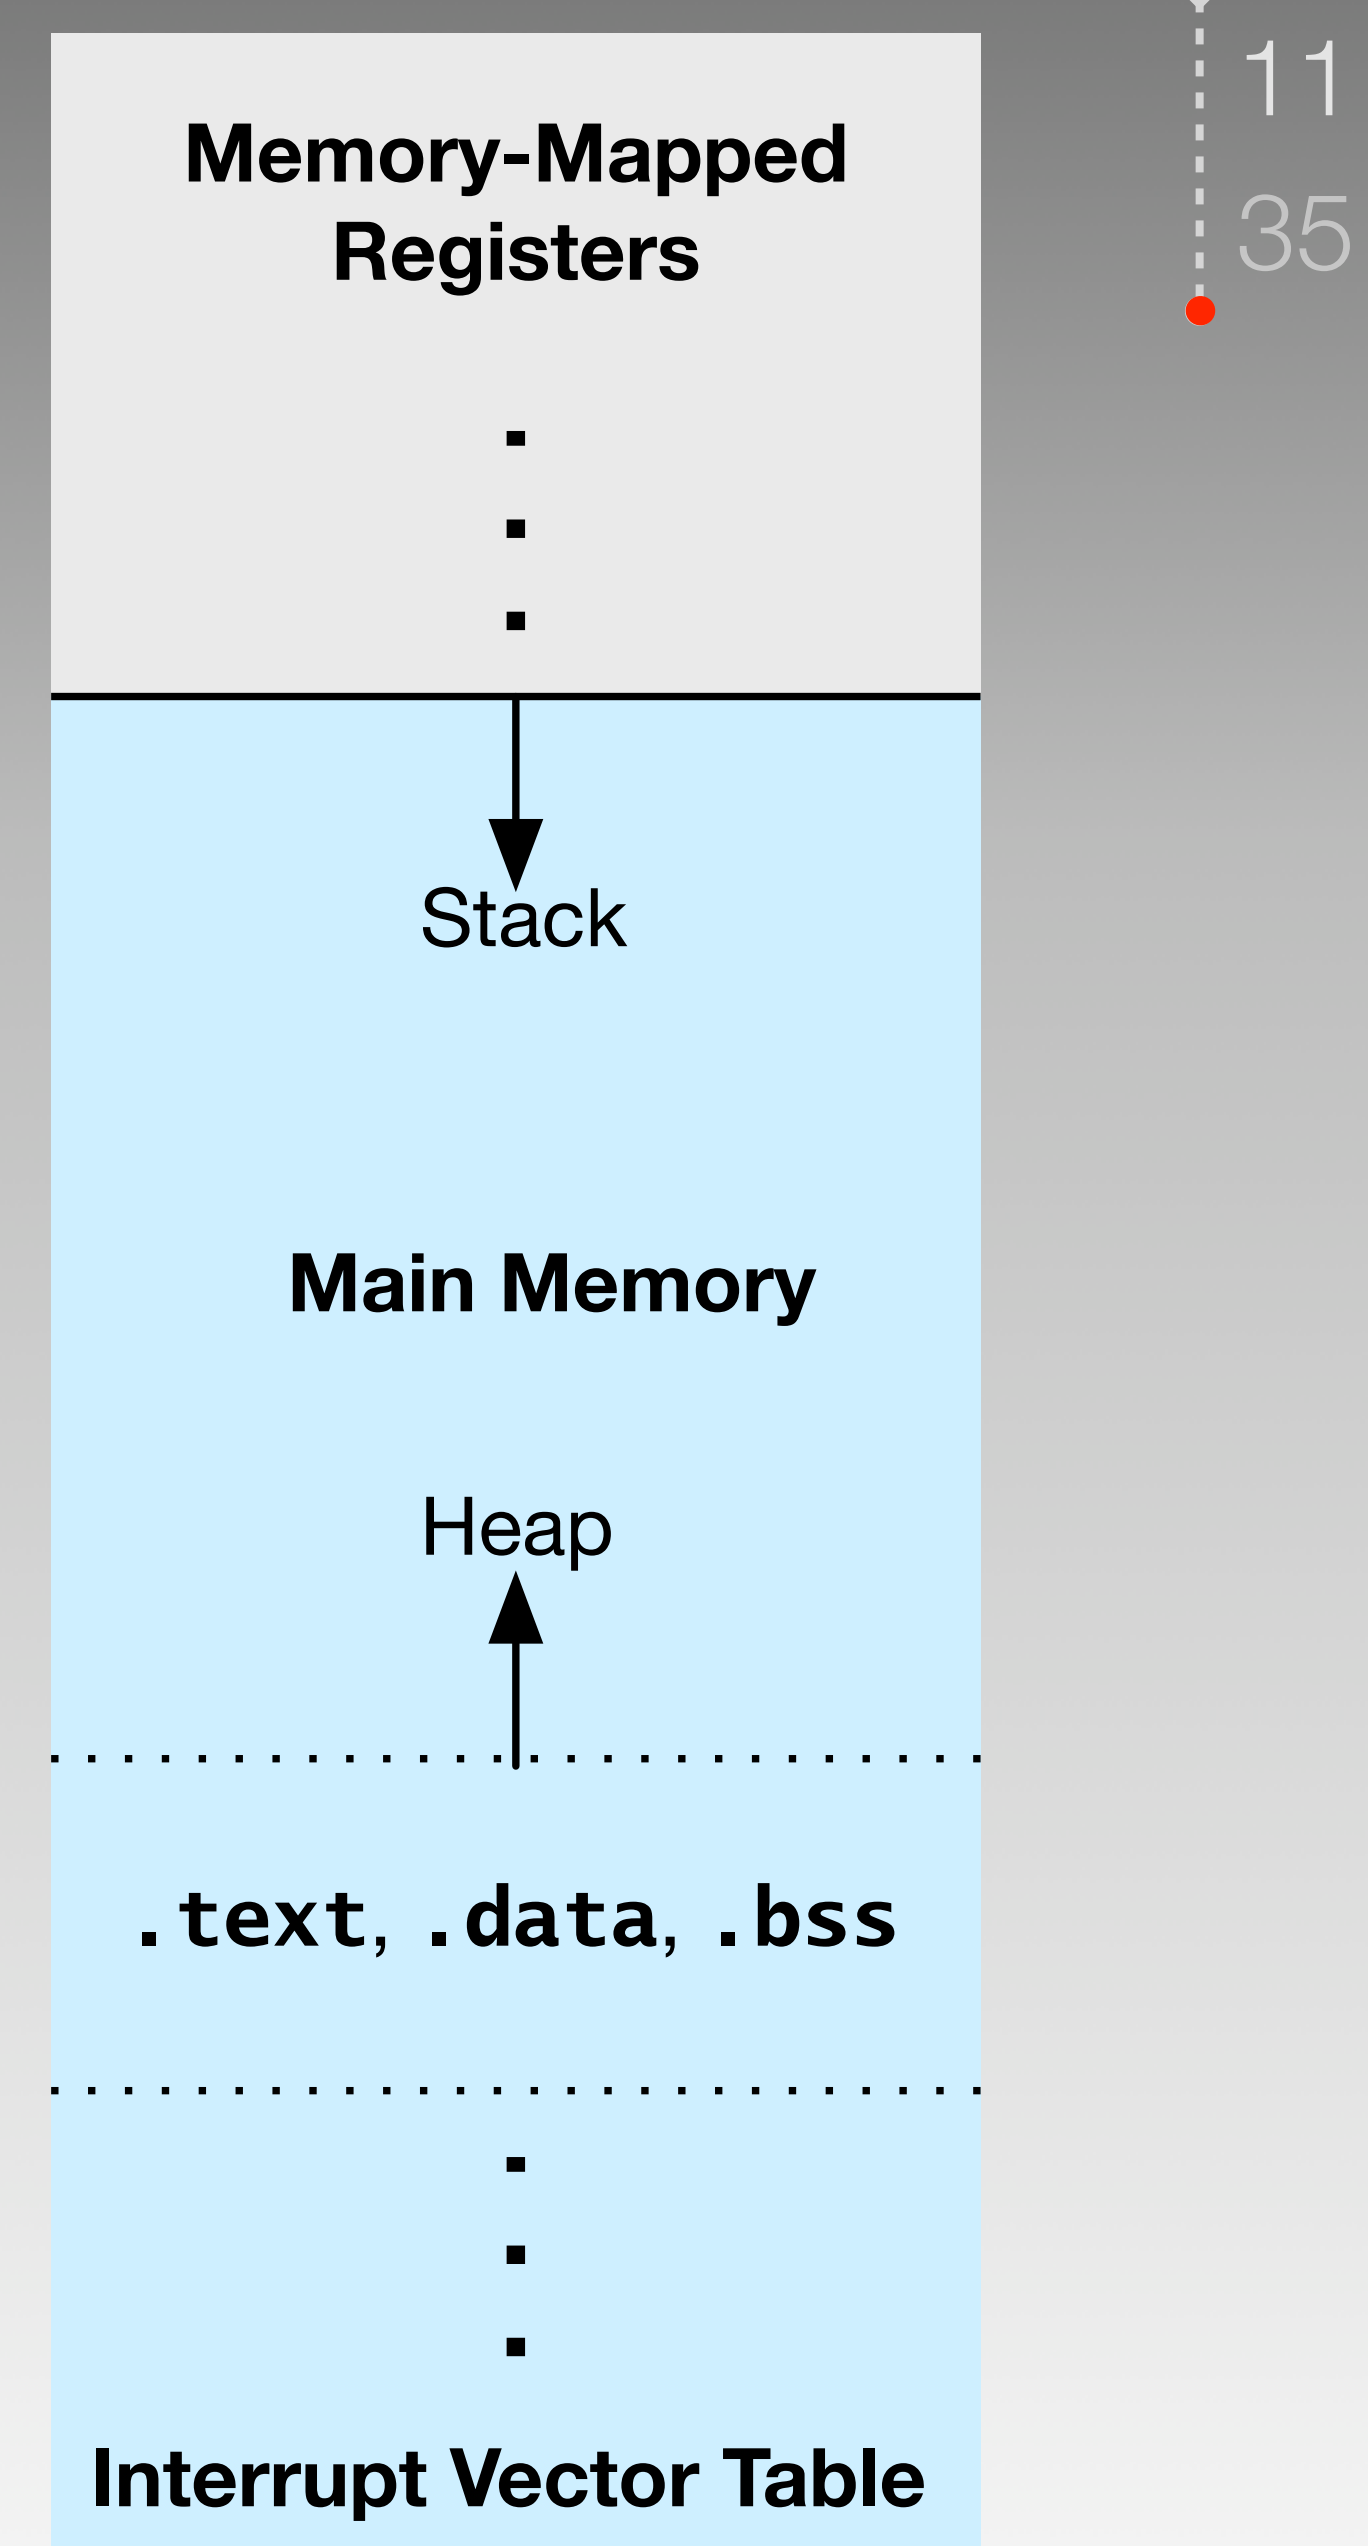
Terminology

- ▶ Memory map
- ▶ Memory-mapped registers
- ▶ Stack
- ▶ Heap
- ▶ **.text** section
- ▶ **.data** section
- ▶ **.bss** section



How do we know what is where in memory/binary?

- ▶ Stack?
- ▶ Heap?
- ▶ **.text?**
- ▶ **.data?**
- ▶ **.bss?**



Like maps in the physical world, **map files** tell us what's where in the binary

Map File:

- ▶ An **output file** generated by the linker which indicates what symbol names are in the final binary, which object files they came from, their sizes, and where in the binary they are located. (Map files often have the filename extension **.map**)

map | mæp |

noun

1 a diagrammatic representation of an area of ...

What's In a Binary? ^{1 of 3}: Understanding Map Files

13

35

```
1 int gInitializedInt = 42;
2 int gUninitializedInt;
3
4 int
5 noMain(void)
6 {
7     int i, localInt;
8
9     for (i = 0; i < 10; i++) {
10         localInt += gInitializedInt + gUninitializedInt;
11     }
12
13     return 0;
14 }
```

What's In a Binary? ^{2 of 3}: Understanding Map Files

```
1 int gInitializedInt = 42;
2 int gUninitializedInt;
3
4 int
5 noMain(void)
6 {
7     int i, localInt;
```

Map File Lines 1 to 7:

```
1 Allocating common symbols
2 Common symbol      size      file
3
4 gUninitializedInt   0x4        simple.o
5
6 Memory Configuration
7
```


What's In a Binary? 3 of 3: Understanding Map Files

15

35

```
1 int gInitializedInt = 42;
2 int gUninitializedInt;
3
4 int
5 noMain(void)
```

Map file, lines 13 to 21:

```
13 Address of section .text set to 0x8004000
```

```
14
15 .text          0x00000000008004000      0x80
16              0x00000000008004000      _text = .
17 *(.text)
18 .text          0x00000000008004000      0x20 init.o
19              0x00000000008004018      __errno
20 .text          0x00000000008004020      0x60 simple.o
21              0x00000000008004020      noMain
22              0x00000000008004080      _etext = .
```

Debugging Information:

- ▶ Typically part of (stored in) binary, but not loaded into processor for execution. **Contains information that allows you to correlate processor state to program source code.**

What's In a Binary? ^{1 of 2}: Debugging Information

17

35

```
1 int gInitializedInt = 42;
2
3 int
4 noMain(void) {
5     int i, localInt;
6
7     for (i = 0; i < 10; i++) {
8         localInt += gInitializedInt + gUninitializedInt;
```

“Stabs” debugging information (dumped), lines 1 to 10:

```
1 % sh-coff-objdump -G simple
2
3 simple:      file format coff-sh
4
5 Contents of .stab section:
6
7 Symnum n_type n_othr n_desc n_value  n_strx String
8
9 -1      HdrSym 0      31      00000000000000292 1
10 0       S0     0       2       00000000008004020 1      simple.c
```


What's In a Binary? 2 of 2: Debugging Information

18

35

```
1 int gInitializedInt = 42;
2
3 int
4 noMain(void) {
5     int i, localInt;
6
7     for (i = 0; i < 10; i++) {
8         localInt += gInitializedInt + gUninitializedInt;
```

“Stabs” debugging information (dumped), lines 27 to 36:

27	17	FUN	0	0	00000000008004020	573	noMain:F(0,1)
28	18	SLINE	0	6	00000000000000000	0	
29	19	SLINE	0	9	00000000000000006	0	
30	20	SLINE	0	11	00000000000000012	0	
31	21	SLINE	0	9	0000000000000002a	0	
32	22	SLINE	0	14	00000000000000044	0	
33	23	SLINE	0	15	00000000000000048	0	
34	24	LSYM	0	0	00000000000000000	587	i:(0,1)
35	25	LSYM	0	0	00000000000000004	595	localInt:(0,1)
36	26	LBRAC	0	0	00000000000000000	0	

What's In a Binary? 4 of 4: Understanding Map Files

19

35

```
1 int gInitializedInt = 42;
2 int gUninitializedInt;
3
4 int
5 noMain(void)
```

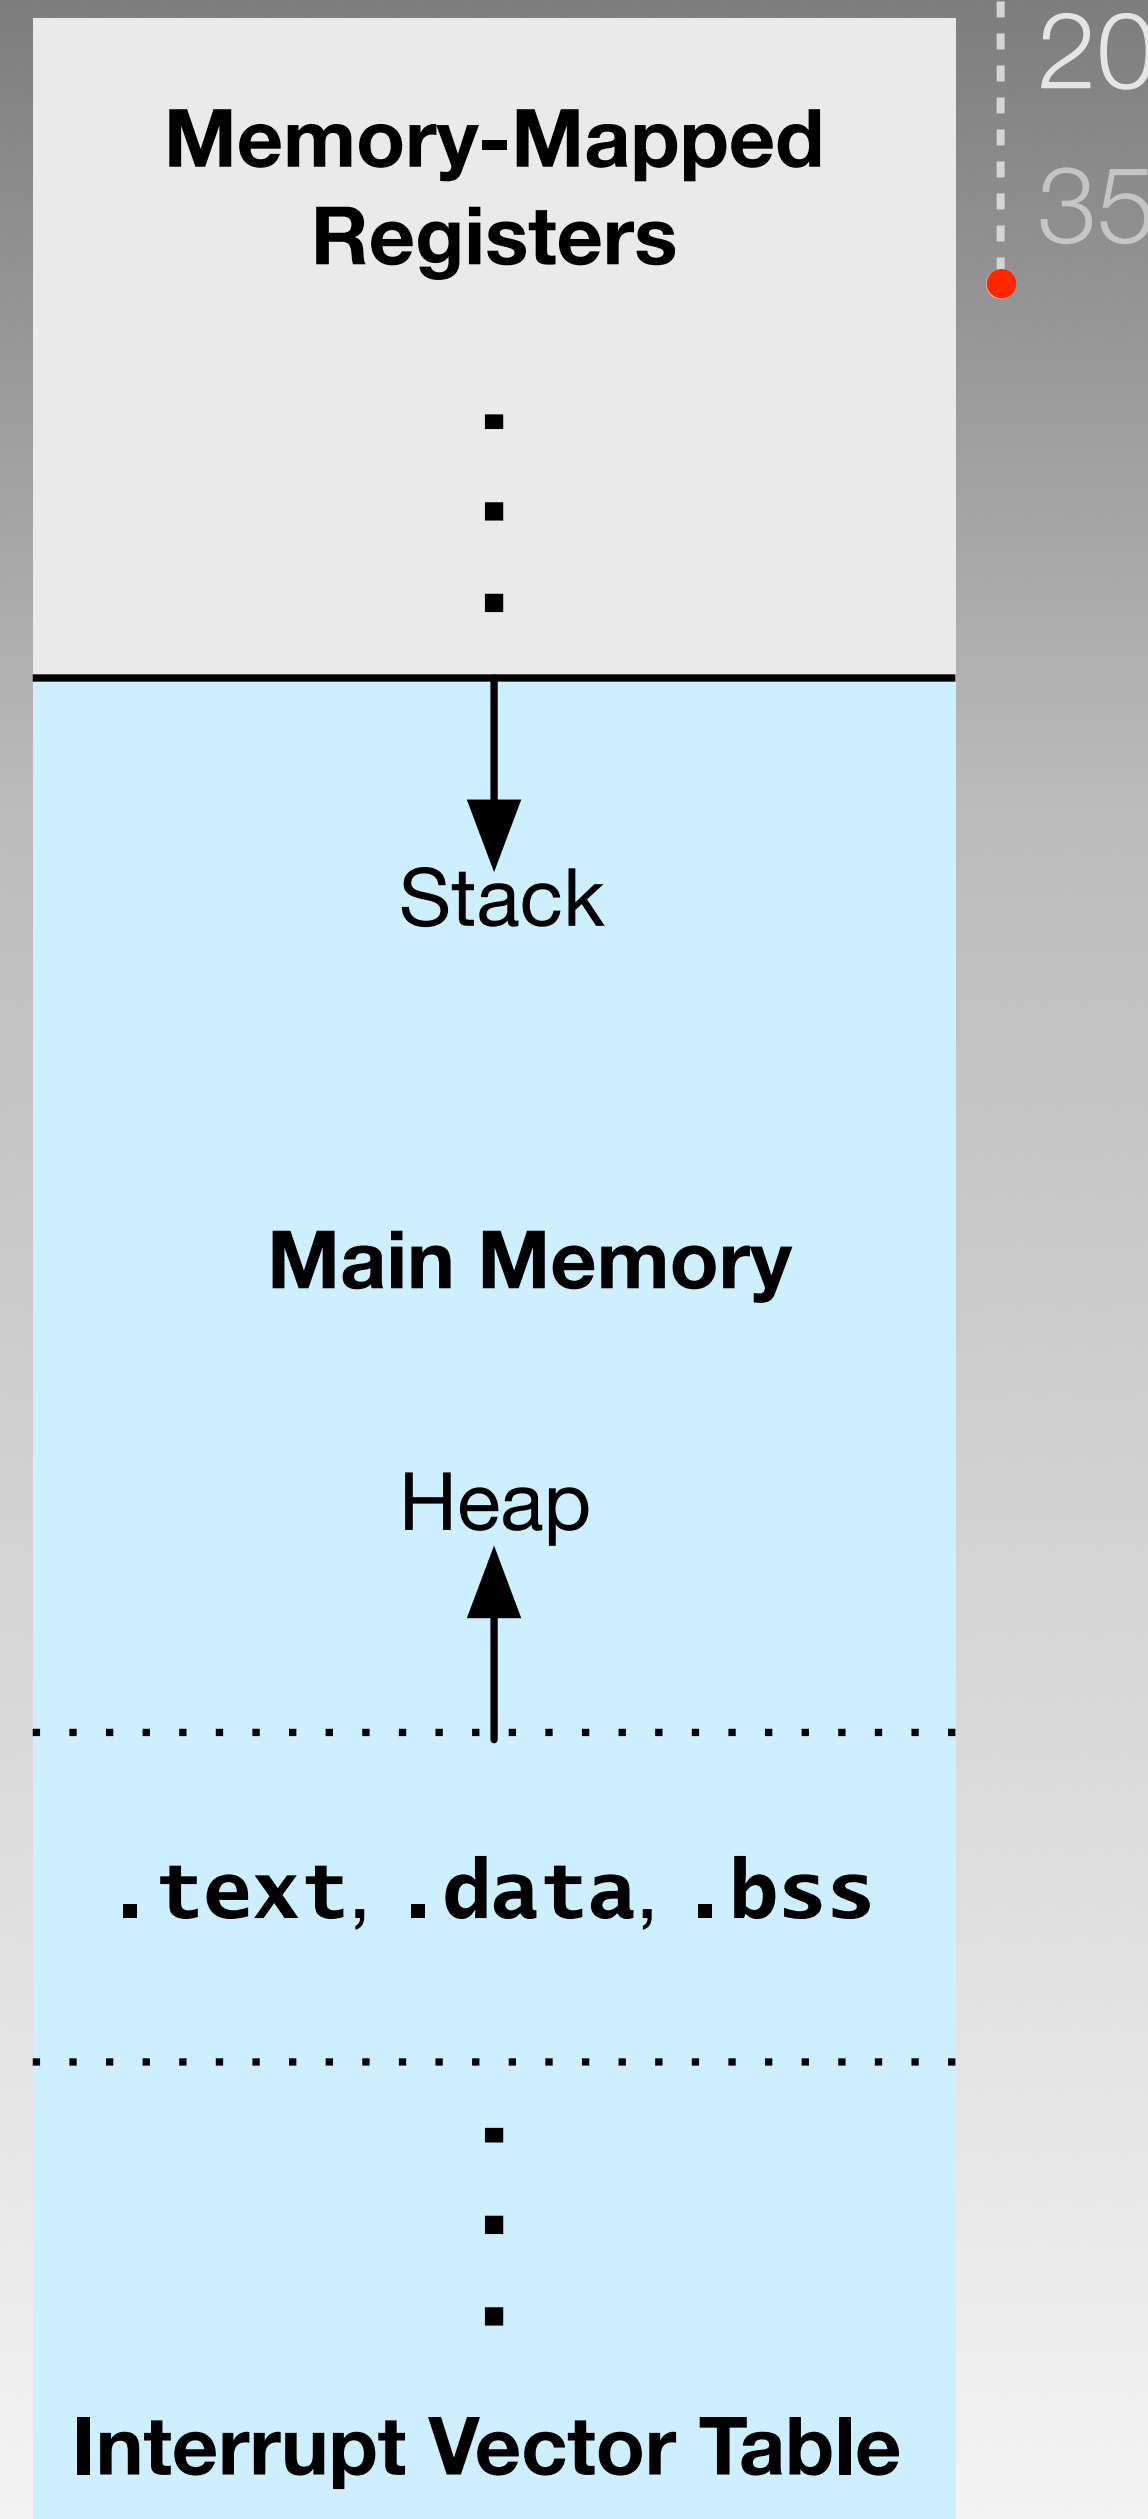
Map File

```
LOAD init.o
LOAD simple.o
LOAD ../../../../tools/tools-lib/superH/libc.a
LOAD ../../../../tools/tools-lib/superH/libgcc.a
LOAD ../../../../tools/tools-lib/superH/libm.a
LOAD ../../../../sys/libOS/mOS/libmOS-superH.a
OUTPUT(simple coff-sh)
```

.stab	←	0x00000000008004094	0x180	
.stab		0x00000000008004094	0x180	simple.o

How do we control what goes where in memory/binary?

- ▶ Stack?
- ▶ Heap?
- ▶ **.text?**
- ▶ **.data?**
- ▶ **.bss?**



Linker command files tell the linker how to layout the contents of binaries

Linker Script File (Linker Command File)

- ▶ An **input command file**, used by the linker, which tells the linker how to place code in the final linked binary.

command | kə'mænd |

noun

an authoritative order: it's unlikely they'll obey your commands.

- Computing an instruction or signal that causes...

Linker Script Files^{1 of 2}

22

35

```
OUTPUT_FORMAT("coff-sh")
OUTPUT_ARCH(sh)
SECTIONS
{
    .text :
    {
        _text = . ;
        *(.text)
        *(.strings)
        _etext = . ;
    }
    .tors :
    {
        __ctors = . ;
        *(.ctors)
        __ctors_end = . ;
        __dtors = . ;
        *(.dtors)
        __dtors_end = . ;
    }
    11 lines deleted...
    .bss :
    {
        _bss = . ;
        *(.bss)
        *(COMMON)
        _ebss = . ;
        _end = . ;
    }
}
```

Linker Script Files^{2 of 2}

23

35

```
OUTPUT_FORMAT("coff-sh")
OUTPUT_ARCH(sh)
SECTIONS
{
```

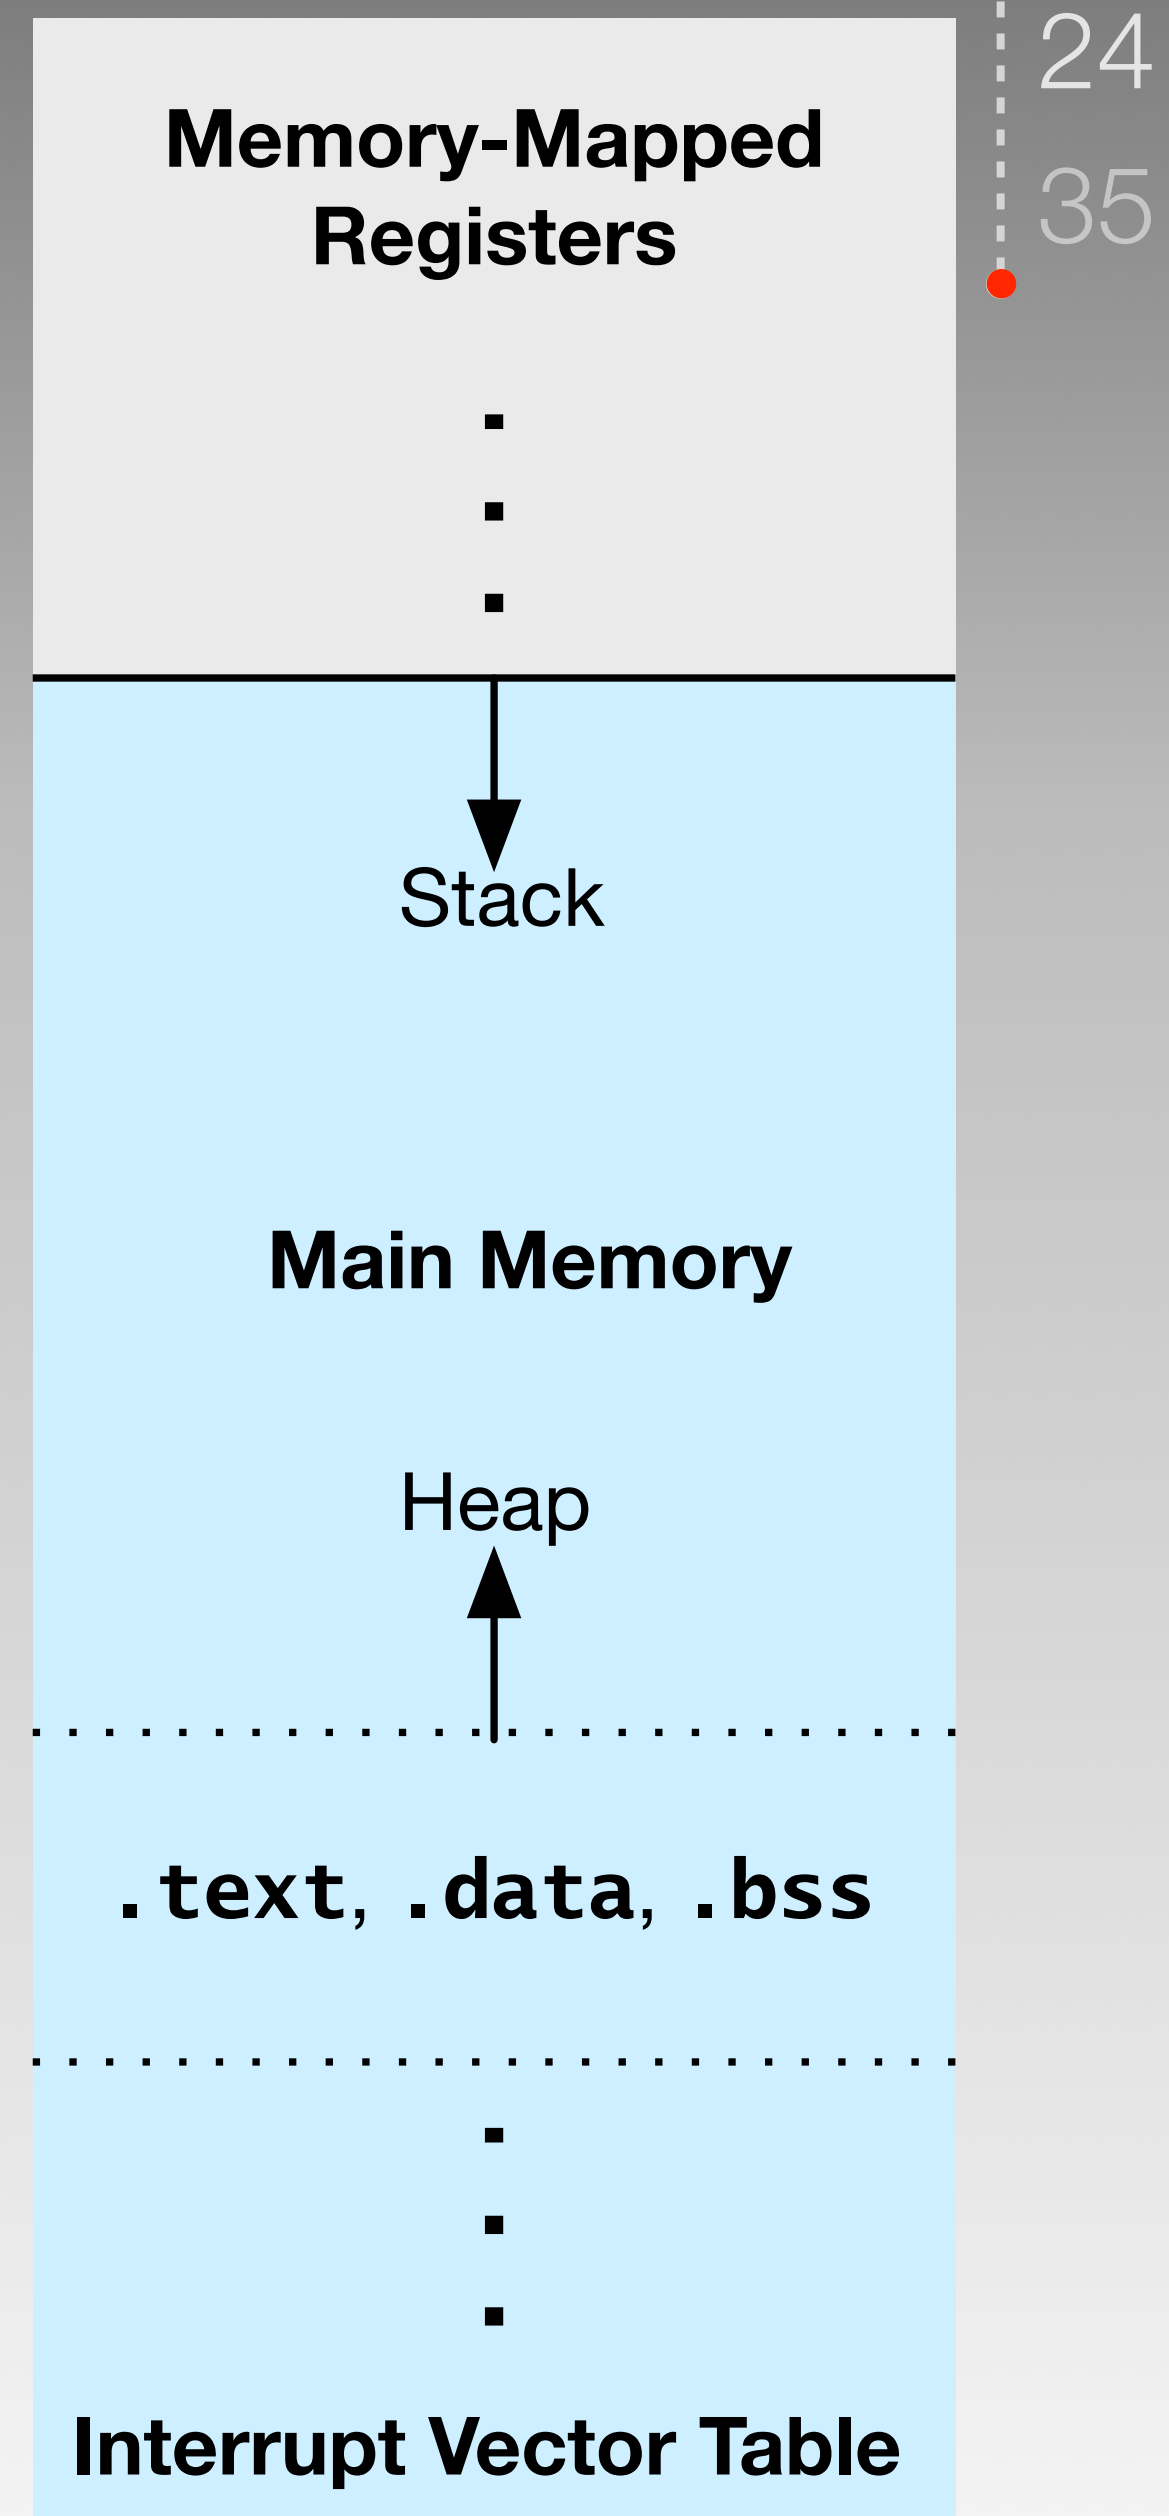
```
    .text :
    {
        _text = . ;
        *(.text)
        *(.strings)
        _etext = . ;
    }
```

```
    .ctors :
    {
        __ctors = . ;
        *(.ctors)
        __ctors_end = . ;
    }
```


Recap

How do we control what goes where in memory/binary?

- ▶ Stack?
- ▶ Heap?
- ▶ **.text?**
- ▶ **.data?**
- ▶ **.bss?**



Linker command files tell the linker how to layout the contents of binaries

Next:

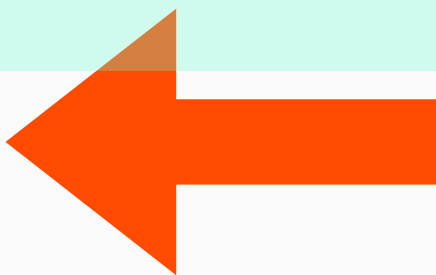
- ▶ We will look at **an extended example that combines C and assembly language** to implement **processor initialization** and a **relocated interrupt handler**

Combining C, Assembler, and an Interrupt Handler^{1 of 9}

26

35

```
1  #include "print.h"
2
3  /*      8 GPRs + PR      */
4  unsigned char  REGSAVESTACK[36];
5  static void    hdlr_install(void);
6  volatile int   gFlag = 1;
7
8  int
9  main(void) {
10     hdlr_install();
11
12     print("\n\n\nStarting...\n\n\n");
13     while (gFlag) {
14     }
15     print("\n\n\nExiting, bye!\n\n\n");
16
17     return 0;
18 }
```



and...


```
19
20 void
21 intr_hdlr(void) {
22     gFlag = 0;
23
24     return;
25 }
26
```


Combining C, Assembler, and an Interrupt Handler ^{2 of 9}

27

35

```
27 void
28 hdlr_install(void) {
29     extern unsigned char    vec_stub_begin, vec_stub_end;
30     unsigned char *         dstptr = (unsigned char *)0x8000600;
31     unsigned char *         srcptr = &vec_stub_begin;
32
33     /* Copy the vector instructions to vector base */
34     while (srcptr < &vec_stub_end)
35         *dstptr++ = *srcptr++;
36 }
37
38 return;
39 }
```



Key Concept: explicitly referencing a specific memory address (e.g., **0x8000600**) in C



Combining C, Assembler, and an Interrupt Handler

3 of 9

28

35

```
1  #include "asm.h"
2
3  .global _vec_stub_begin
4  .global _vec_stub_end
5  .global _sleep
6  .global _atexit
7
8  .align 2
9  start:
10     /* Clear Status Reg */
11     AND #0, r0
12     LDC r0, sr
13
14     /* Go ! */
15     MOVL    app_stack_addr, r15
16     MOVL    start_addr, r0
17     JSR @r0
18     NOP
```

Note 1

Note 2

F(E)



Combining C, Assembler, and an Interrupt Handler^{4 of 9}

29

35

```
28      .align 4
29  _vec_stub_begin:
30      /*      Save PR      */
Note 3 MOVL    savestack_addr, r0
32      ADD #36, r0
33      STS.L   pr, @-r0
34
35      /*      Save R8 - R15    */
36      BSR saveregs
37      NOP
38
39      /*  It's now safe to call  */
40      MOVL    intr_stack_addr, r15
```

Recall

```
27  void
28  hdlr_install(void) {
29      extern unsigned char  vec_stub_begin, vec_stub_end;
30      unsigned char *      dstptr = (unsigned char *)vec_stub_end;
31      unsigned char *      srcptr = &vec_stub_begin;
32
33      /*  Copy the vector instructions to vector stub  */
34      while (srcptr < &vec_stub_end) {
35          *dstptr++ = *srcptr++;
36      }
37
38      return;
39  }
```

Key Concept: Symbols (e.g., **vec_stub_begin**) in assembly/C/linker script file



Combining C, Assembler, and an Interrupt Handler 5 of 9

```
41      MOVL      hdlr_addr, r0
42      JSR @r0
43      NOP
44
45      /* Restore R8 - R15      */
46      BSR restorerregs
47      NOP
48
49      /* Restore PR          */
Note 3 MOVL      savestack_addr, r0
51      ADD #32, r0
52      LDS.L    @r0+, pr
53
54      /* Return from exception */
55      RTE
56      NOP
```

Combining C, Assembler, and an Interrupt Handler

6 of 9

31

35

```
59      /*                                     */
60      /*  SR.RB == 1. Save R8-R15. We store items in reverse */
61      /*  so that we can use MOV.L Rm, @-Rn, rather than have */
62      /*  to incr the memory store address ourselves.      */
63      /*                                     */
64  saveregs:
65      /*  Addr of bottom of save area */
66      MOV.L    savestack_addr, r0
67
68      /*  Get addr _end_ of stack */
69      ADD #32, r0
70
71      /*  Now store items bkwrds */
72      MOV.L    r15, @-r0
73      MOV.L    r14, @-r0
74      MOV.L    r13, @-r0
```

Note 3

MOV.L savestack_addr, r0

F(E)

(10)



Combining C, Assembler, and an Interrupt Handler^{7 of 9}

32

35

```
85      /*                                     */
86      /*      SR.RB == 1.  Restore R8-R15.      */
87      /*                                     */
88  restorerregs:
89      /*  Addr of bottom of save area  */
Note 3  MOVL    savestack_addr, r0
91
92      /*  Pop into approp. reg  */
93      MOVL    @r0+, r8
94      MOVL    @r0+, r9
95      MOVL    @r0+, r10
96      MOVL    @r0+, r11
97      MOVL    @r0+, r12
98      MOVL    @r0+, r13
```


Combining C, Assembler, and an Interrupt Handler^{8 of 9}

33

35

```
104
105     .align 4
106     hdlr_addr:
107     .long    _intr_hdlr
108
109     .align 4
110     savestack_addr:
111     .long    _REGSAVESTACK
112
113     .align 4
114     intr_stack_addr:
115     .long    (0x80000000 + (1 << 16))
116     _vec_stub_end:
```

Note 3!!!

Key Concept: array **REGSAVESTACK[]** in the C code provides a place to stash saved state



Combining C, Assembler, and an Interrupt Handler^{9 of 9}

34

35

```
125
126     .align 2
127     /* Stack is 32k above us */
128 app_stack_addr:
129     .long (0x80000000 + (1 << 15))
130
131     .align 2
132 start_addr:
133     .long _main
```

Key Concept: `intr_hdlr()` and `main()` have separate 32k stacks. **32k big enough?** KL03?

Further Reading

35

36

Best next step: Get some practice

- ▶ Like learning to swim, you can't learn all you need from a textbook

A basic C programming tutorial you might find useful:

<https://www.edx.org/course/programming-in-c-getting-started>

Test your understanding:

- ▶ Complete these online self-assessments on <https://f-of-e.org/>
<https://f-of-e.org/chapter-04/#exercises>

Things to Do

36

36

- 1 Complete a “**muddiest point**” 2-question survey using this [link](#)

$F(E)$

